

OUTLINE

- Images:
 - Intensity histogram
 - Noise
- Image filtering:
 - Linear filter
 - Low-pass filtering
 - Non-linear filter
- 2D Fourier Transform

INTENSITY HISTOGRAM

The histogram is a graph showing the number of pixels in an image at each different intensity value found in that image.

Histogram $h(r_k) = n_k$

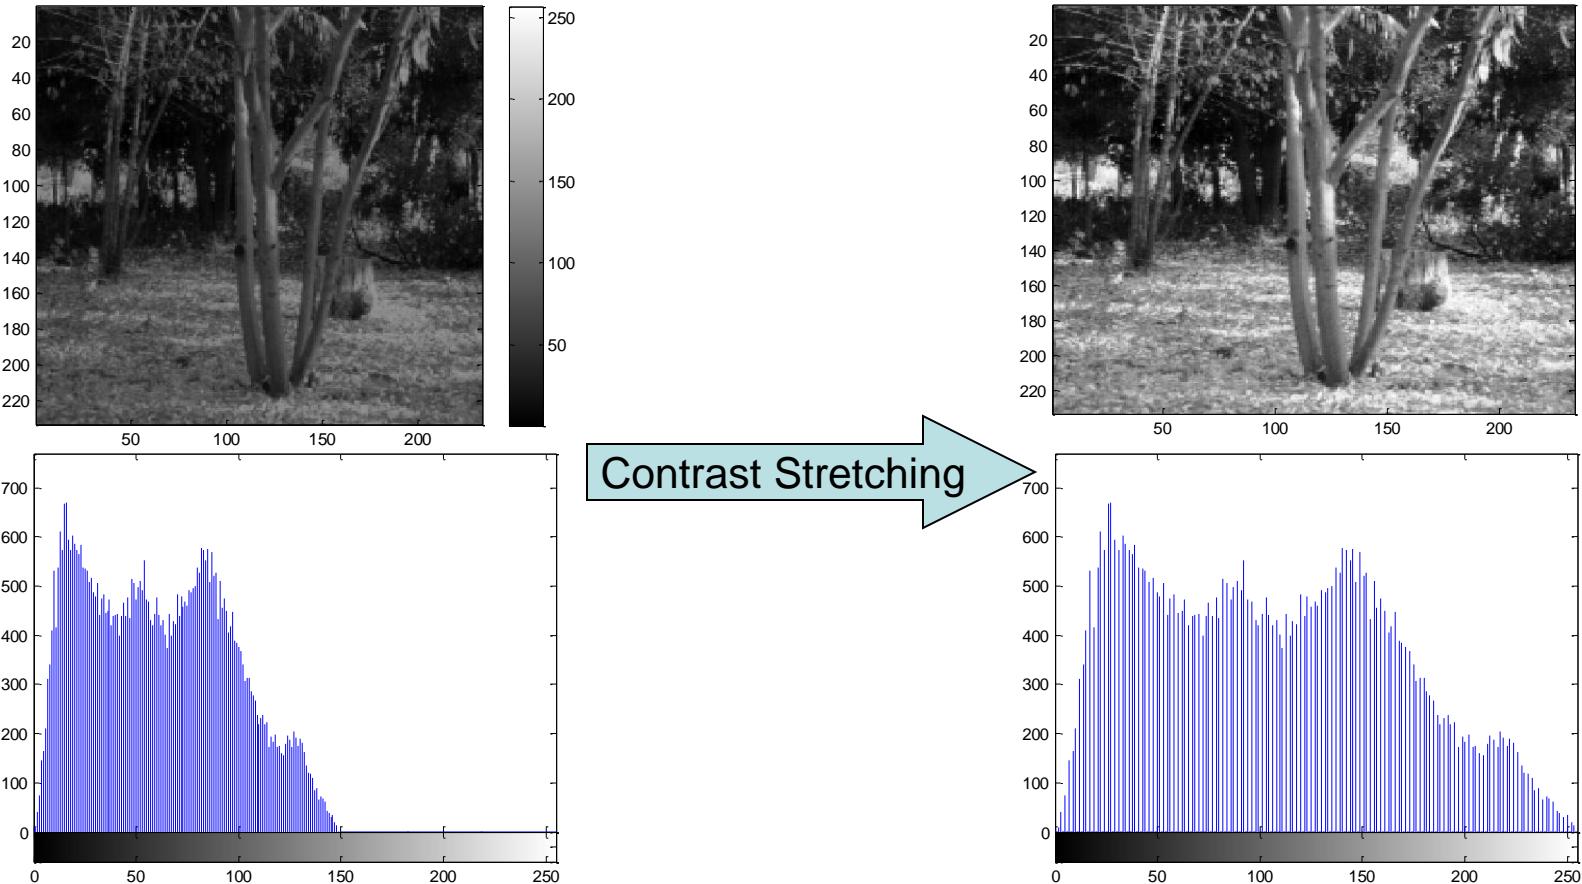
r_k is the k^{th} intensity value

n_k is the number of pixels in the image with intensity r_k

Contrast Stretching: expands the range (a_{min} , a_{max}) of intensity levels in an image $a(i,j)$ so that it spans the full intensity range ($omin$, $omax$) of the recording medium or display device.

$$b(i,j) = ((a(i,j) - a_{min}) / (a_{max} - a_{min})) (omax - omin) + omin$$

INTENSITY HISTOGRAM



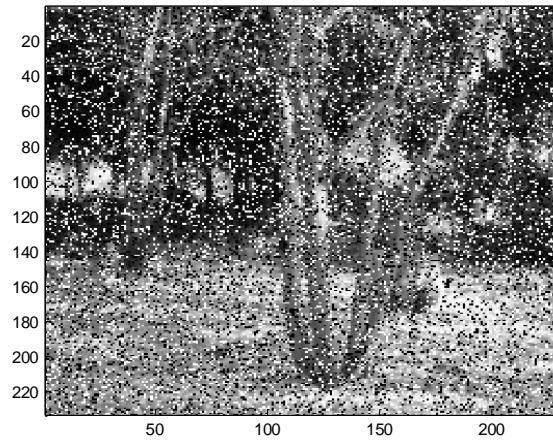
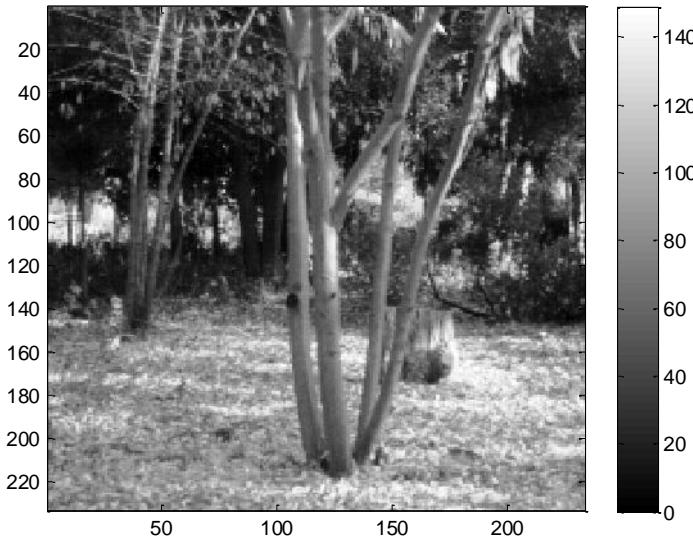
NOISE

- Digital images are corrupted by noise during image formation process (e.g. light fluctuations, sensor noise, quantization effects, finite precision).
- Common types of noise (we often assume the noise is additive):
 - $I(x,y) = s(x,y) + n_i$
 - Where $s(x,y)$ is the deterministic signal
 - n_i is a random variable
 - Gaussian noise: variations in intensity drawn from a Gaussian normal distribution.

NOISE

- Impulse (“shot”) noise, i.e. salt and pepper noise: contains random occurrences of black and white pixels.
- Digital images can be corrupted artificially in order to assess the performance of various vision processing algorithms.

NOISE



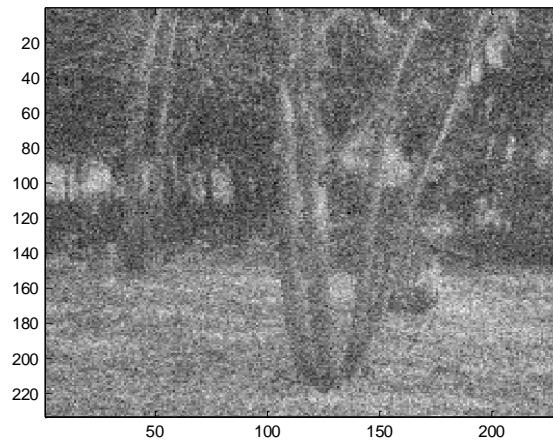
Noise density = 0.3

Gaussian
noise



$\sigma=10$

Computer Vision



$\sigma=30$

See noise.m

IMAGE PROCESSING

- Standard image processing operators map pixel values from one image to another:
- *Point operators*: where each output pixel's value only depends on the corresponding input pixel value (e.g., contrast adjustments and color transformations).
- *Neighborhood (area-based) operators*: where each new pixel's value depends on a small number of neighboring input pixel values (e.g. filtering).
- *Global operators* (e.g. histogram and Fourier Transform).

IMAGE (linear) FILTERING

- Form new image whose pixels are a weighted sum of original pixel values, using the same set of weights at each point of the image.
- Output is a linear function of the input.
- Output is a shift-invariant function of the input (i.e. shift the input image two pixels to the left, the output is shifted two pixels to the left).

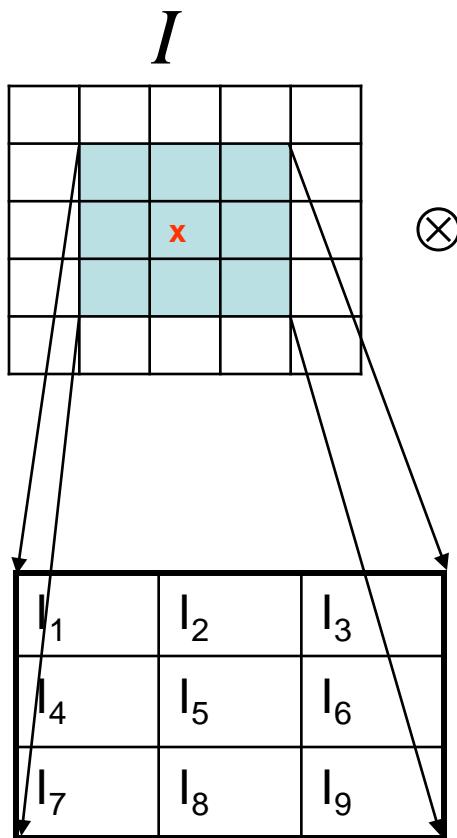
IMAGE FILTERING

- In order to form a new image whose pixels are a weighted sum of original pixel values, using the same set of weights (some function of a local neighborhood of the pixel) at each point:

$$O(i, j) = I * H = \sum_k \sum_l I(k, l)H(i - k, j - l)$$

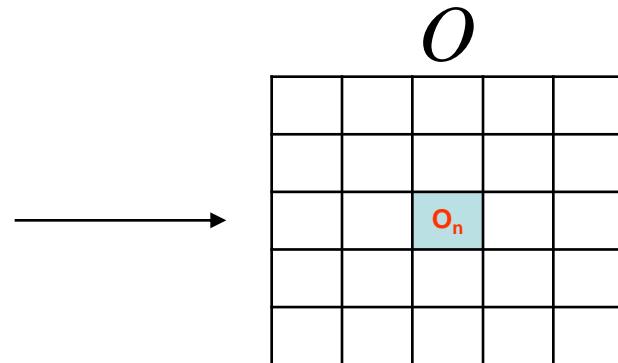
- where I is the input image, O is the output image and H is the convolution kernel (the two-dimensional spatial filter).

IMAGE FILTERING



$$I \otimes H = I_1 H_1 + I_2 H_2 + I_3 H_3 + I_4 H_4 + I_5 H_5 + I_6 H_6 + I_7 H_7 + I_8 H_8 + I_9 H_9$$

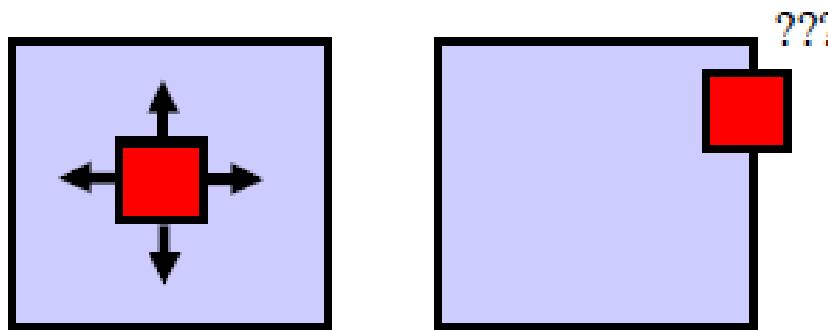
We don't want to only do this at a single pixel, of course, but want instead to "run the kernel over the whole image".



$$\begin{aligned}
 O_n = & I_1 H_9 + I_2 H_8 + I_3 H_7 \\
 & + I_4 H_6 + I_5 H_5 + I_6 H_4 \\
 & + I_7 H_3 + I_8 H_2 + I_9 H_1
 \end{aligned}$$

IMAGE FILTERING

- Problem: what do we do for border pixels where the kernel does not completely overlap the image?



- Different border handling methods specify different ways of defining values for pixels that are off the image.
- One of the simplest methods is **zero-padding**.

IMAGE FILTERING

image

Origin $f(x, y)$

0 0 0 0 0

0 0 0 0 0

0 0 1 0 0

0 0 0 0 0

0 0 0 0 0

kernel

$w(x, y)$

1 2 3

4 5 6

7 8 9

Padded f

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

Rotated w

9 8 7

6 5 4

3 2 1

0 0 0

0 0 0

0 0 0

0 0 0

0 0 0

0 0 0

Full convolution result

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 1 2 3 0 0 0 0

0 0 0 4 5 6 0 0 0 0

0 0 0 7 8 9 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

0 0 0 0 0 0 0 0 0 0

Cropped convolution result

0 0 0 0 0

0 1 2 3 0

0 4 5 6 0

0 7 8 9 0

0 0 0 0 0

SMOOTHING

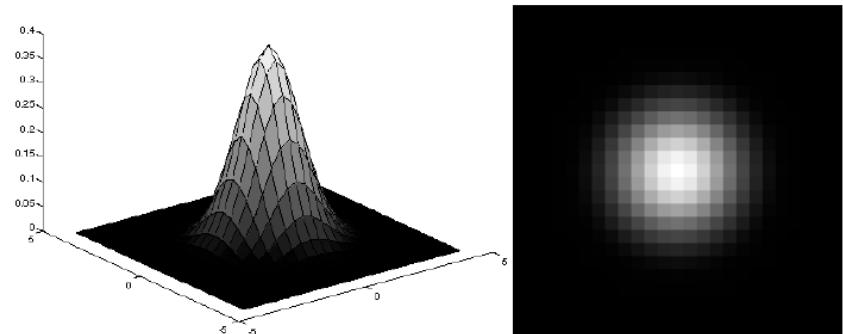
- Smoothing:
 - Smoothing filters are used for blurring and for noise reduction
 - Blurring is used in removal of small details and bridging of small gaps in lines or curves
 - Smoothing spatial filters include linear filters and nonlinear filters.
- Smoothing (linear filters):
 - by averaging (box filter) performs the average of pixels in a neighborhood;
 - with a Gaussian (low-pass filter) performs a weighted average of pixels in a neighborhood.

SMOOTHING

Box filter

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

since this is a linear operator, we can take the average around each pixel by convolving the image with this 3x3 filter

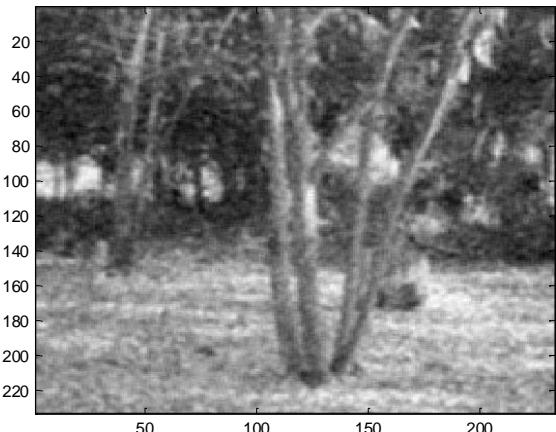
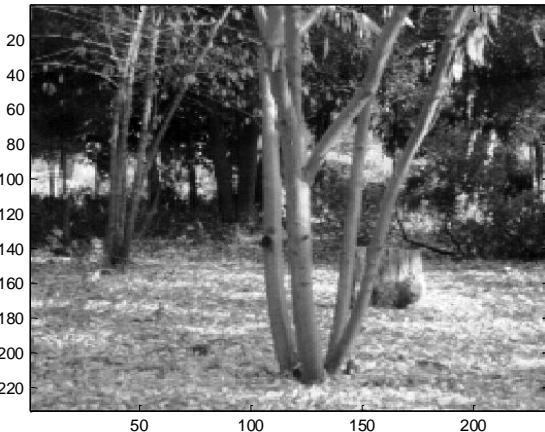


Gaussian Smoothing Filter (weighted averaging): the coefficients are a 2D Gaussian, i.e. it gives more weight at the central pixels and less weights to the neighbors

$$G_{\sigma} \equiv \frac{1}{2\pi\sigma^2} \exp\left\{-\frac{x^2+y^2}{2\sigma^2}\right\}$$

IMAGE FILTERING: SMOOTHING

- Smoothing reduces noise (for appropriate noise models): since the pixels “are like” their neighbors, whereas the noise is independent from pixel to pixel.



$\sigma=20$

average

IMAGE FILTERING: Efficient Implementation

- Both, the Box filter and the Gaussian filter are separable:
 - First convolve each row with a 1D filter
 - Then convolve each column with a 1D filter.

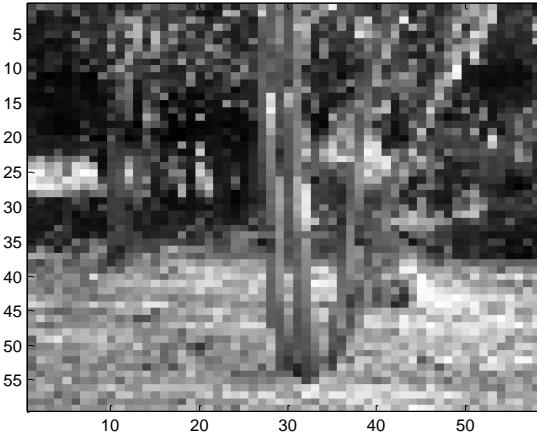
Separable Gaussian: **associativity**

$$G_\sigma * f = [g_{\sigma \rightarrow} * g_{\sigma \uparrow}] * f = g_{\sigma \rightarrow} * [g_{\sigma \uparrow} * f]$$

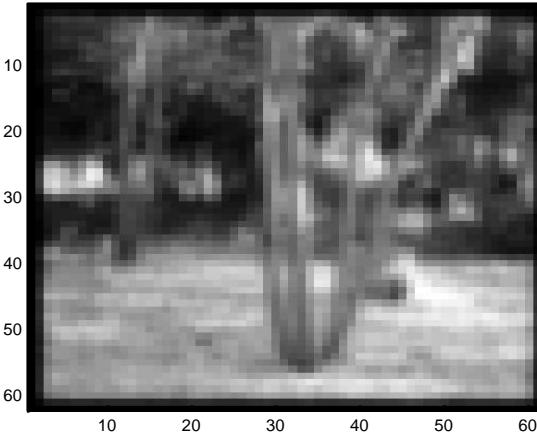
IMAGE RESOLUTION: SUBSAMPLING



See subsampling.m



Sub-sampling



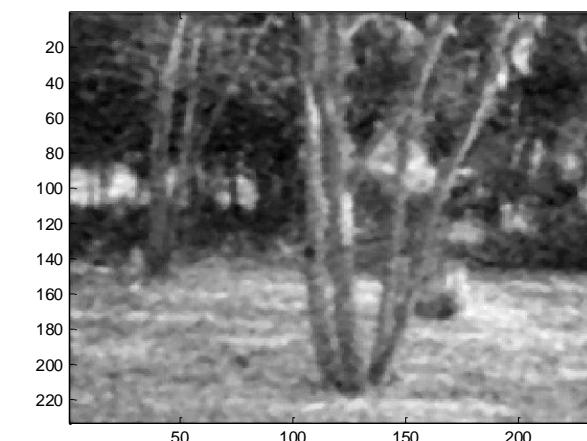
Low-pass filtering and sub-sampling

MEDIAN FILTER

- It is a non-linear filter: (1) rank-order neighborhood intensities and (2) take middle value.
 - median completely discards the spikes;
 - median preserves discontinuities;
 - median loses important details and produces patchy effect.



See noise_filtering.m



average

median

2D FOURIER TRANSFORM

- We define the Fourier transform of an image $g(x,y)$ to be:

$$\mathfrak{F}(g(x,y))(u,v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} g(x,y) e^{-i2\pi(ux+vy)} dx dy$$

where the result is a complex valued function of (u, v) .

- Any function that is not periodic can be expressed as the integral of sines and /or cosines multiplied by a weighing function
- The transform of an digital image is performed by FFT.

2D FOURIER TRANSFORM: 2D impulse

The impulse $\delta(x, y)$,

$$\delta(x, y) = \begin{cases} 1 & \text{if } x = y = 0 \\ 0 & \text{otherwise} \end{cases}$$

The sifting property

$$\sum_{x=-\infty}^{\infty} \sum_{y=-\infty}^{\infty} f(x, y) \delta(x, y) = f(0, 0)$$

and

$$\sum_{x=-\infty}^{\infty} \sum_{y=-\infty}^{\infty} f(x, y) \delta(x - x_0, y - y_0) = f(x_0, y_0)$$

2D FOURIER TRANSFORM: properties

Translation

$$f(x, y)e^{j2\pi(\mu_0x/M + \nu_0y/N)} \Leftrightarrow F(\mu - \mu_0, \nu - \nu_0)$$

and

$$f(x - x_0, y - y_0) \Leftrightarrow F(\mu, \nu)e^{-j2\pi(\mu x_0/M + \nu y_0/N)}$$

Convolution

$$f(x, y) * h(x, y) \Leftrightarrow F(u, v)H(u, v)$$

$$f(x, y)h(x, y) \Leftrightarrow F(u, v) * H(u, v)$$

2D FOURIER TRANSFORM: properties

Fourier Spectrum and Phase Angle

2-D DFT in polar form

$$F(u, v) = |F(u, v)| e^{j\phi(u, v)}$$

Fourier spectrum

$$|F(u, v)| = \left[R^2(u, v) + I^2(u, v) \right]^{1/2}$$

Power spectrum

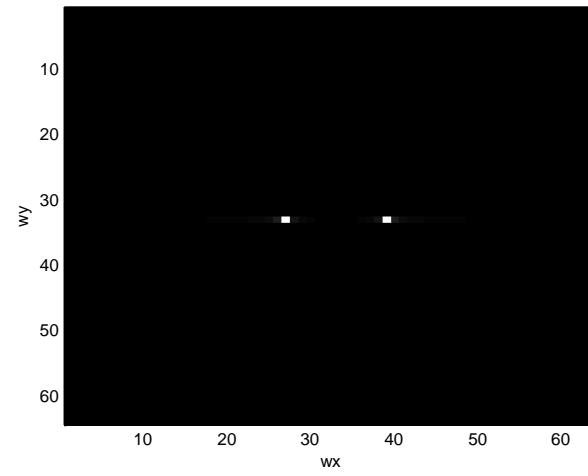
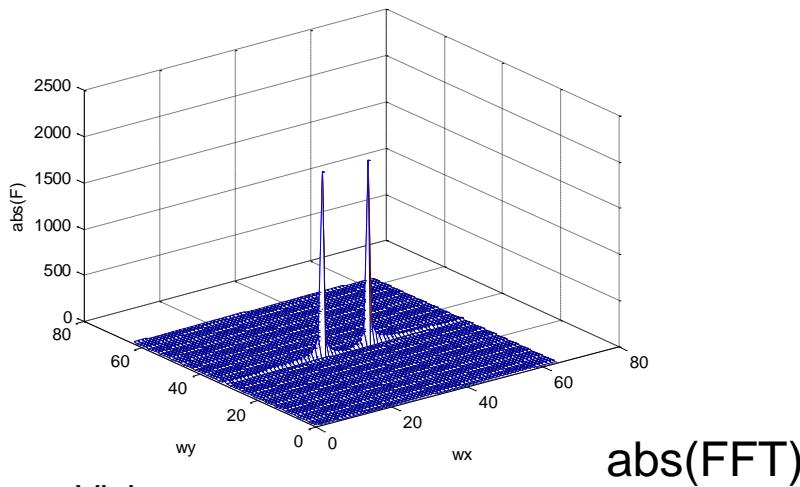
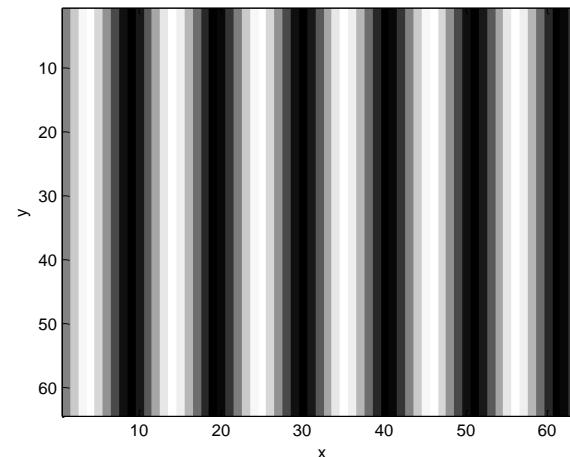
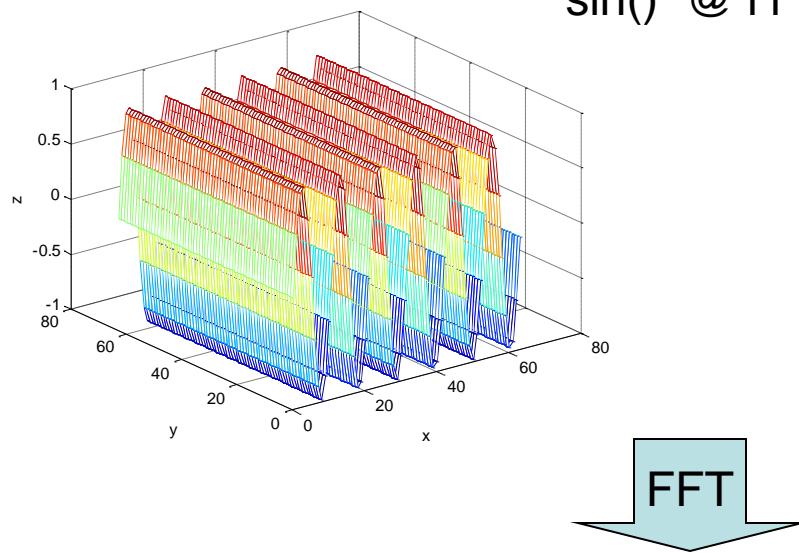
$$P(u, v) = |F(u, v)|^2 = R^2(u, v) + I^2(u, v)$$

Phase angle

$$\phi(u, v) = \arctan \left[\frac{I(u, v)}{R(u, v)} \right]$$

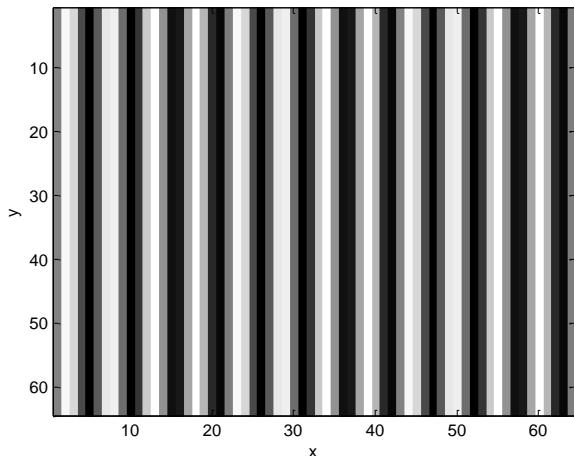
2D FOURIER TRANSFORM

$\sin()$ @ f_1

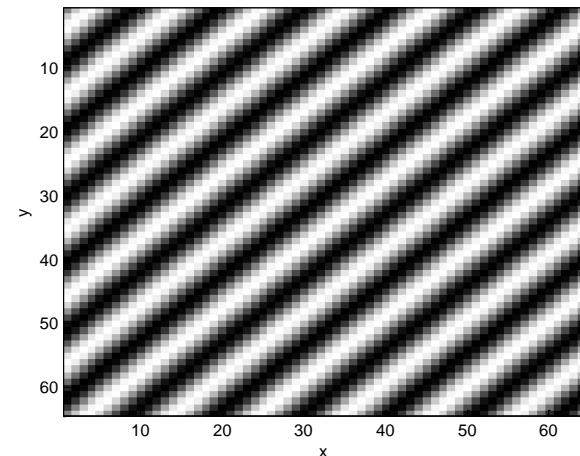


2D FOURIER TRANSFORM

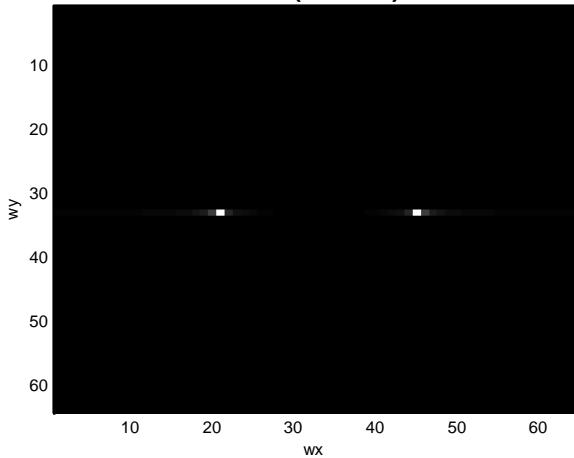
$\sin()$ @ $f_2 > f_1$



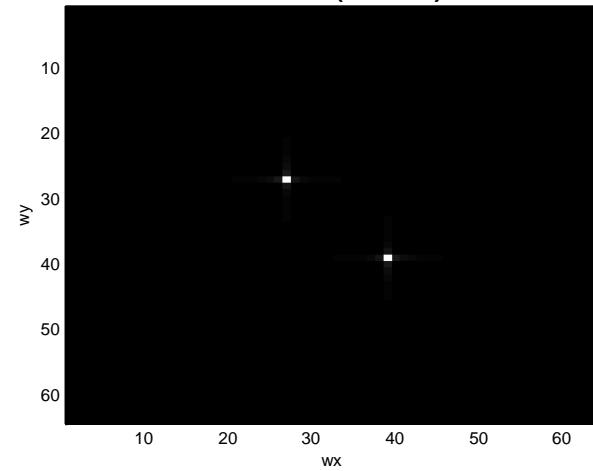
$\sin()$ @ f_1 , different orientation



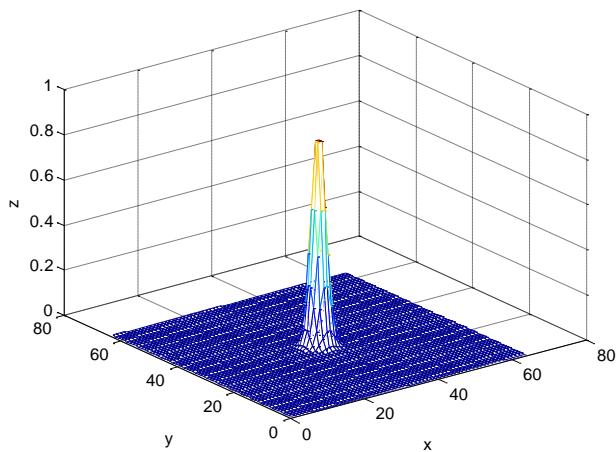
abs(FFT)



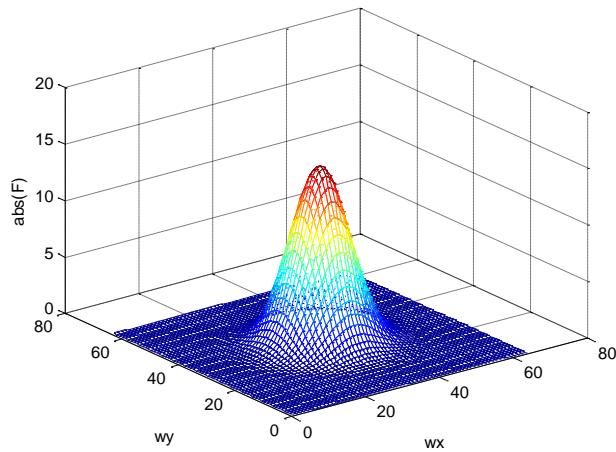
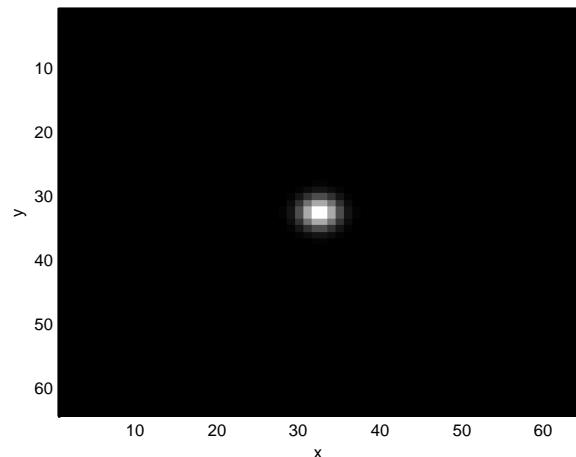
abs(FFT)



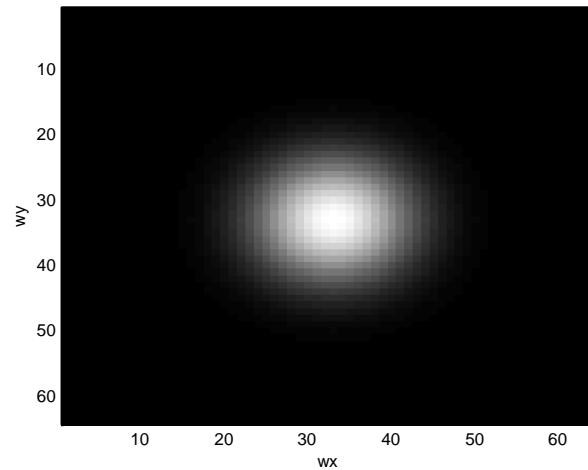
2D FOURIER TRANSFORM



Gaussian

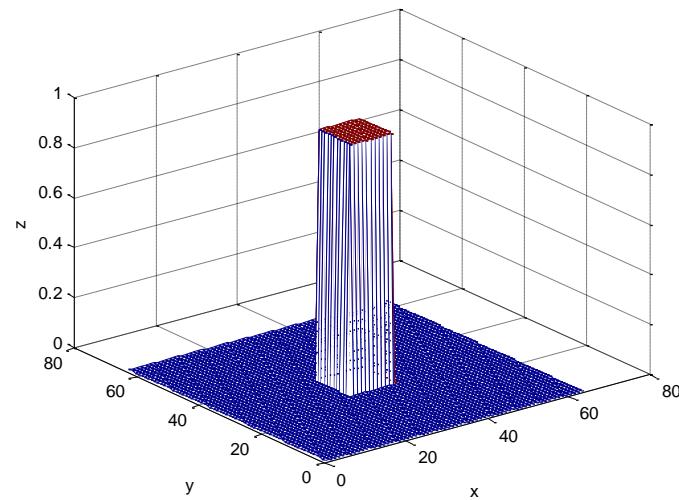


abs(FFT)

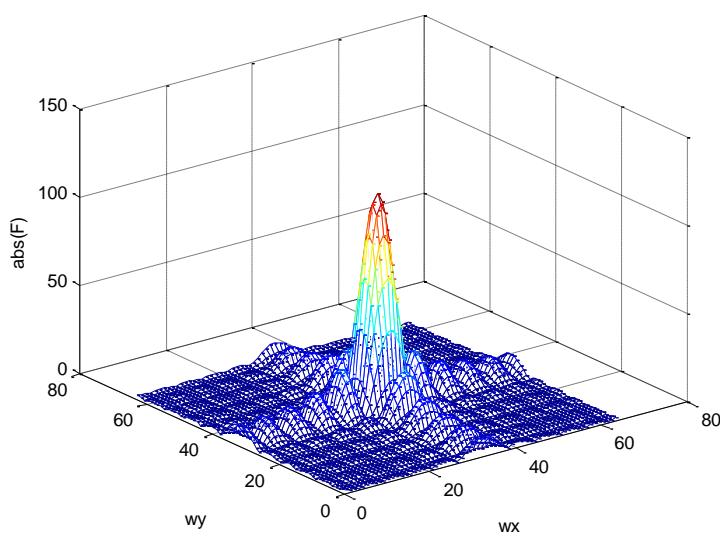
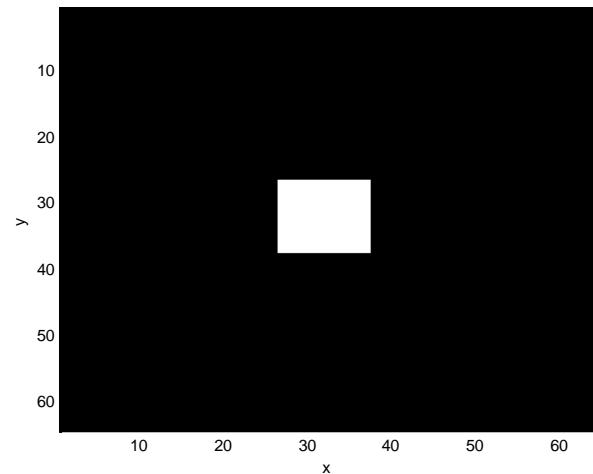


Low-pass filter

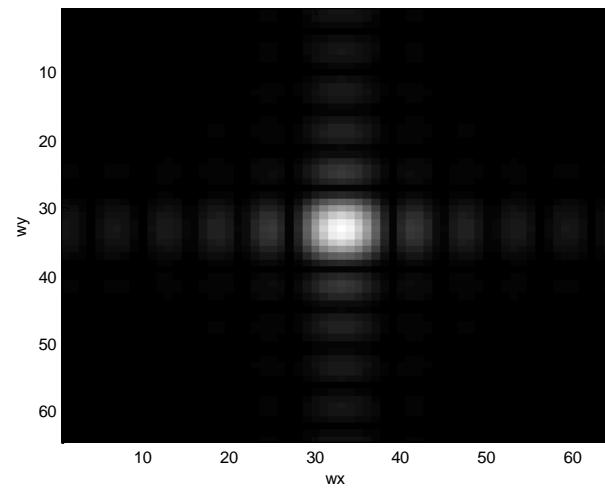
2D FOURIER TRANSFORM



rectangle



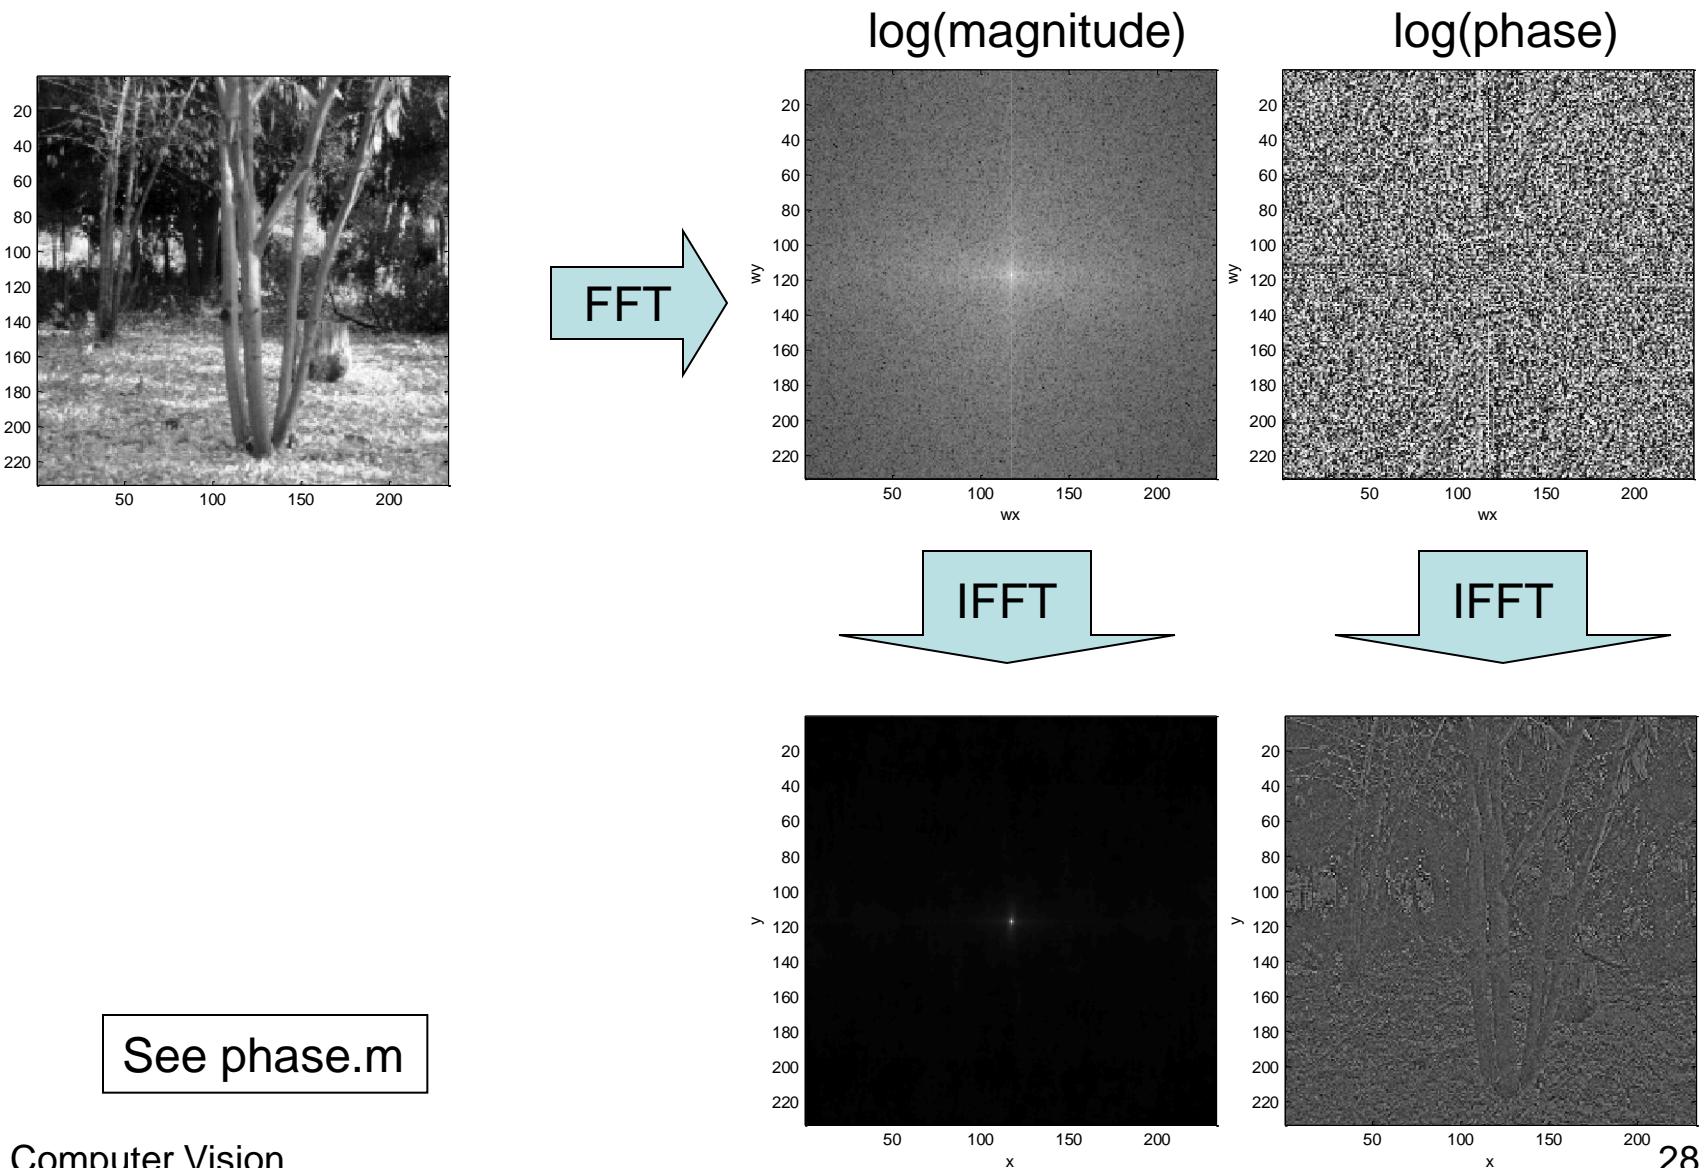
$\text{abs}(FFT)$



2D FOURIER TRANSFORM

- The Fourier transform consists of a real and a complex component: the magnitude and phase, respectively.
- The value of the Fourier transform of a function/image at a particular (u,v) point depends on the whole function/image.
- The phase information is crucial to reconstruct the correct spatial structure of the image.

2D FOURIER TRANSFORM



HOMEWORK

- Reading Assignments:
 - Mubarak Shah, “Fundamentals of Computer Vision”: sections 2.1, 2.2, 2.3, 2.4, 2.5, 2.6, 2.10.
- Lab : Friday, October 3, 11.00-13.00 (E5)